

Modeling Profit Give-Back Risk in Algorithmic Trading via LSTM-Based Exit Decision Systems

First Author*

Second Author[†]

This Draft: June 1, 2026

Abstract

This paper examines the profit give-back problem in trend-following algorithmic trading caused by the inherent lag of rule-based exit signals. We propose a two-layer decision framework that separates market entry and exit tasks. While a baseline Weighted Moving Average (WMA) crossover strategy determines position entry, a Long Short-Term Memory (LSTM) network is trained to optimize the timing of early liquidation. In addition to standard market features, the model incorporates position-context variables, including maximum favorable excursion (MFE) and trailing drawdown, to capture transaction-level risks. Using high-frequency cryptocurrency data (BTC, ETH, XRP, and SOL) from 2022 to 2026, we empirically evaluate the out-of-sample performance of the framework. The results indicate that the LSTM-driven exit system significantly improves cumulative returns and Sharpe ratios relative to both the baseline WMA strategy and a passive buy-and-hold benchmark. Furthermore, granular per-trade PnL distribution analysis shows that the framework effectively mitigates downside drawdown risks while preserving the returns generated by macroscopic market trends. **Keywords:** LSTM Model; Cryptocurrency; Profit Give-Back Risk; Exit Strategy Optimization.

JEL Classification: E43; E44; E52; G12; G15.

Acknowledgments: We thank ... for helpful comments and suggestions. All errors are our own.

*First Author: Department of Finance, XXX University. Email: first.author@email.com.

[†]Second Author: School of Economics, XXX University. Email: second.author@email.com.

1 Introduction

Problem formulating: In quantitative trading, trend-following strategies especially moving average (MA) or weighted moving average (WMA) crossover strategies are widely used due to their structural simplicity and cross-market robustness. These strategies rely on MA crossover signals to initiate and reverse positions, effectively capturing medium-to long-term price trends. While entry rules have been extensively studied and well-established, exit timing optimization remains a critical bottleneck limiting strategy performance.

Traditional trend-following approaches typically exit positions only when an opposite crossover signal occurs, a mechanism that suffers from inherent lag. Although it captures full trend cycles, it often fails to lock in accumulated unrealized profits during late-stage trends or range-bound markets, leading to profit drawdowns. Existing solutions, such as fixed take-profit/stop-loss rules or trailing stops, rely on empirically calibrated parameters, struggle to adapt to dynamic market conditions, and lack data-driven frameworks for early exit decisions.

To systematically characterize this problem, we conduct an empirical analysis of the trading paths of traditional WMA strategies in the cryptocurrency market and identify a salient empirical fact that has not been sufficiently discussed: the vast majority of trades once achieved positive unrealized returns during the holding period, but only a small proportion eventually closed with positive returns. Specifically, in the dataset examined in this study, around 90% of trades were in profit at some point, whereas only about 30% ended as profitable trades (see in Table 1). Whereas the profit path (see in Figure 1) shows a highly consistent dynamic structure: in the early stage of the holding period, profits increase monotonically over time; in the middle-to-late stage, profits reach a peak corresponding to MFE; and in the later stage, a significant drawdown occurs. This finding suggests that the main bottleneck of trend-following strategies is not the absence of profit opportunities, but the lack of an effective mechanism for profit locking and risk control.

Table 1: Statistical Results of Trading Paths

Asset	Total Trades	MFE > 0 Trades	MFE > 0 Ratio	Final PnL > 0 Trades	Final PnL > 0 Ratio	Average MFE (USDT)	Average Final PnL (USDT)	Total PnL (USDT)
DOGE	1575	1371	87.05	566	35.94	36.75	2.79	4389.06
SOL	1579	1371	86.83	551	34.90	38.79	2.60	4103.44
XRP	1579	1359	86.07	552	34.96	30.58	1.77	2789.82
ETH	1607	1404	87.37	541	33.67	26.04	1.34	2152.53
BTC	1627	1399	85.99	536	32.94	19.02	0.91	1486.19

Note: MFE refers to Maximum Favorable Excursion, and PnL denotes Profit and Loss. All statistical results are quoted in USDT, and all ratio indicators are presented as percentages (%). The 30-minute frequency cryptocurrency trading data utilised in this paper is retrieved from the public API of Binance, spanning the sample period from January 1, 2022 to 2026. Within the adopted WMA strategy, the lengths of the long and short trading k-line windows are set to 60 and 30, respectively.




Figure 1: Average Profit Paths Across Cryptocurrencies

Note: This figure illustrates the average trade profit corresponding to each cryptocurrency asset. The horizontal axis denotes the holding duration, while the vertical axis represents profit measured in USDT. The shaded area corresponds to the 95% confidence interval.

Based on the above results, this paper formalizes the phenomenon as unrealized profit drawdown risk and models it as a sequential decision-making problem: during the holding period, how can one determine, based on the current market state and position-related characteristics, whether continued holding will lead to a significant future drawdown in returns, and thereby decide whether the position should be exited early. Unlike traditional approaches that focus on price prediction or trend identification, this paper emphasizes the dynamic evolution of the profit path and its associated risk structure.

We propose an LSTM-based early-exit decision framework that reformulates the exit problem as a sequence classification task. The framework is designed to be embedded as a data-driven decision layer atop traditional trend-following strategies, learning the optimal exit boundary by balancing future returns and potential drawdowns. Backtesting on cryptocurrency market data demonstrates that the proposed method significantly improves key performance metrics while overall returns remain stable.

The remainder is organized as follows. Section 2 reviews the related literature in this field.

The remainder of the paper is organized as follows. Section 2 reviews the related literature in this field. Section 3 describes the dataset used in this study and how we split the sample. Section 4 displays our strategy and how it works. Section 5 reports the results of both the baseline strategy and ours, and shows that our LSTM early stopping model outperforms the baseline strategy. Section ?? uses both expanding windows and rolling windows to re-execute the experiment to ensure the robustness of our strategy. Section 7 concludes the paper and points out future research directions.

2 Related Work

Quantitative trading research has long focused on identifying price regularities and transforming them into implementable strategies. The literature has evolved from traditional technical analysis to machine learning-based prediction and, more recently, to reinforcement learning-based sequential decision making. While extensive studies have explored entry signals and price forecasting, optimal exit timing—especially for mitigating unrealized profit drawdown—remains understudied.

2.1 Trend-following Strategies

Trend-following represents one of the most classical paradigms in quantitative trading, which exploits price inertia through technical indicators to identify and trade in line with prevailing trends. Moving-average-based methods are the most representative implementations, including the simple moving average (SMA), weighted moving average (WMA), exponential moving average (EMA), and short–long moving-average crossover rules. Hunter (1986) provided an early systematic discussion of the exponentially weighted moving average, emphasizing the importance of recent information in time-series analysis. Hansun (2013) further discussed the evolution and improvement of traditional moving-average tools.

As research developed, scholars began to examine the effectiveness of moving-average rules from theoretical and empirical perspectives rather than treating them as purely empirical tools. Chiarella et al. (2006) analyzed the mechanism of moving-average rules

from a dynamic market perspective, noting that window length and trading feedback significantly affect price dynamics. [Ellis and Parbery \(2005\)](#) compared adaptive and simple moving-average strategies, showing that dynamic parameter adjustment improves adaptability. [Marshall et al. \(2017\)](#) compared time-series momentum and moving-average rules, highlighting differences in signal construction and return sources.

Overall, trend-following strategies have evolved from fixed-parameter rules toward adaptive and diversified structures. Traditional technical analysis has not become obsolete with the rise of machine learning; instead, it provides important features for intelligent trading models. For instance, [Ayala et al. \(2021\)](#) showed that technical strategies are increasingly optimized and integrated via machine learning, enabling a transition from rule-driven to data-driven frameworks.

2.2 Machine Learning in Trading

The application of machine learning has significantly transformed the paradigm of quantitative trading. Unlike traditional methods that rely on predefined rules, machine learning automatically captures nonlinear mappings between market data and trading outcomes. Early studies such as [Trippi and DeSieno \(1992\)](#) applied neural networks to index futures trading, while [Dempster and Jones \(2001\)](#) and [Dempster and Leemans \(2006\)](#) introduced genetic programming and reinforcement learning, reflecting a shift from manual rule design to automated decision making.

Supervised learning in trading is primarily divided into classification and regression. Classification methods predict price directions or trading signals [Shen et al. \(2012\)](#); [Gerlein et al. \(2016\)](#); [Zhong and Enke \(2019\)](#), whereas regression methods focus on forecasting price levels and returns [Adegboye and Kampouridis \(2021\)](#).

LSTM has become a milestone in financial time-series modeling due to its ability to capture long-term dependencies. [Nelson et al. \(2017\)](#) and [Fischer and Krauss \(2018\)](#) demonstrated its strong performance in direction prediction. [Borovkova and Tsiamas \(2019\)](#) extended LSTM to high-frequency data, and [Mehtab et al. \(2020\)](#); [Ghosh et al. \(2022\)](#) further developed hybrid structures combining LSTM with other models.

Recently, machine learning research has shifted from pure prediction accuracy to strategy-oriented performance. [Sebastião and Godinho \(2021\)](#) highlighted that in-sample precision does not guarantee trading profits. [Han et al. \(2023\)](#) linked label design more closely with trading system construction, indicating a trend toward profit-driven and strategy-aware modeling.

2.3 Exit Strategy Optimization

Compared with entry signal research, exit strategies directly determine profit realization and risk control but have received far less attention. Exit decisions affect not only

trade-level PnL but also drawdown management and risk exposure.

Early exit research focused on fixed take-profit and stop-loss rules. [Leung and Li \(2015\)](#) incorporated stop-loss constraints into mean-reversion strategies, emphasizing the impact of exit boundaries on optimality. [Lo and Remorov \(2017\)](#) showed that stop-loss performance depends heavily on market states. [Vezeris et al. \(2018\)](#) found that dynamic volatility-based thresholds outperform fixed ones.

Trailing stop-loss rules further enable profit locking by adjusting exit levels with favorable price movements. [Glynn and Iglehart \(1995\)](#) provided early theoretical analysis, while [Dai et al. \(2021\)](#) and [Leung and Zhang \(2021\)](#) demonstrated their ability to reduce downside risk and formalized trailing stops as optimal stopping problems.

Reinforcement learning has recently emerged as a tool for sequential exit decisions. [Kim and Kim \(2019\)](#), [Tsantekidis et al. \(2021\)](#), and [Wang et al. \(2021\)](#) applied deep reinforcement learning to optimize trading boundaries and exit quality. However, such methods remain more focused on overall strategy optimization rather than fine-grained exit timing for individual trades.

Despite progress, research on exit optimization remains limited. [Hwang et al. \(2023\)](#) noted that exit rules strongly influence strategy efficiency, while [Sivri et al. \(2023\)](#) and [Shawky \(2023\)](#) represent early attempts to use machine learning for stop-loss and take-profit calibration. The field still lacks systematic, data-driven frameworks for dynamic exit and drawdown control.

2.4 Research Gaps

Existing literature has achieved substantial progress in trend identification, price prediction, and automated trading. However, critical limitations remain in exit optimization, particularly for controlling unrealized profit drawdowns.

First, traditional trend-following strategies rely on fixed parameters and lack robustness to volatile or structurally changing markets. Second, exit research remains underdeveloped relative to entry signal modeling. Studies on trailing exits, dynamic take-profit mechanisms, and integrated exit optimization are especially limited. Third, most studies treat entry, holding, and exit as separate components rather than building a unified decision framework. Poor exit design can erode accurate predictions through profit drawdowns, unnecessary stop-losses, and excessive transaction costs.

In summary, while existing research provides a solid foundation for trend following and machine learning-based prediction, data-driven dynamic exit optimization—especially for profit protection and drawdown control—remains an underexplored area with clear theoretical and practical value.

3 Data Description and Preparation

3.1 Dataset

To objectively evaluate the generalization ability of the LSTM-based early-closing decision framework in dynamic market environments, this study selects the cryptocurrency market as the empirical setting. The cryptocurrency market features 24/7 continuous trading, high volatility, and pronounced trend persistence, making it an ideal setting for testing whether the enhanced strategy can effectively improve the traditional trend-following model.

This study selects four mainstream cryptocurrency trading pairs as the experimental assets: BTC/USDT, ETH/USDT, SOL/USDT, and XRP/USDT, which we collect from the Binance public API. These assets cover different market capitalizations, liquidity characteristics, and price fluctuation patterns, thereby allowing the robustness of the model across assets to be evaluated.

The raw market data used in the experiment range from January 1, 2022 to March 12, 2026. This interval fully covers multiple bull-bear cycles in the market, enabling the model to be tested across different market regimes, including trend continuation, range-bound fluctuation, and sharp drawdown. In this paper, 30-minute candlestick data are used for modeling. This time granularity ensures the reliability of trend signals while providing sufficient sequence depth for the LSTM model, so as to capture the dynamic evolution of micro risk structures during the holding period.

We strictly follows the chronological order principle for predictive samples. Based on the trading id we divides the all the tradings triggered by the baseline strategy into the training set, validation set, and testing set, with proportions of approximately 70%, 15%, and 15%, respectively. The specific sample sizes and corresponding time intervals are shown in Table 2.

Table 2: Dataset Split

Stage	Proportion	Sample Size	Time Interval	Research Function
Training Set	70%	29,250	2022/01/02– 2025/01/06	Model parameter learning and weight optimization
Validation Set	15%	6,268	2025/01/06– 2025/08/05	Hyperparameter tuning and overfitting prevention
Testing Set	15%	6,269	2025/08/05– 2026/03/11	Out-of-sample evaluation and robustness test

Note: Dataset partitioning is based on the total number of trades, with the boundaries strictly defined as $n_{\text{train}} = \lfloor n_{\text{trades}} \times 0.70 \rfloor$ and $n_{\text{val}} = \lfloor n_{\text{trades}} \times 0.15 \rfloor$. Due to the floor truncation operation, a residual fraction of at most one trade is fully allocated to the test set, rendering its actual proportion slightly higher than the theoretical 15%. Given the large sample size, this deviation accounts for less than 0.5% of the total data and exerts a negligible impact on evaluation statistics without compromising the integrity of the train/validation boundaries.

3.2 Data Cleaning

This paper first conducts basic cleaning on the raw candlestick data, including uniformly parsing the time field into UTC time format, converting fields such as prices and trading volume into numerical types, and sorting the data in chronological order. Subsequently, anomaly auditing is performed on the raw sequence to identify missing candlesticks, zero-volume candlesticks, and extreme fluctuation samples, with corresponding flags constructed for each type of anomaly. Since the LSTM model adopts a fixed-length historical window for modeling, we further constructs a context-contamination flag using a rolling-window method, and removes samples whose historical windows contain bad data from the training dataset.

4 Experimental Design

This paper proposes a deep hybrid decision framework that integrates a traditional trend-following strategy with the risk control capabilities of a deep learning model. The framework is composed of two core components: a baseline layer and a Long Short-Term Memory (LSTM) decision layer. The system operates via a two-layer collaborative mechanism. Specifically, the baseline layer utilizes weighted moving average (WMA) crossover signals¹ the LSTM decision layer acts as an intelligent risk management module, monitoring market dynamics to estimate the conditional probability of executing an early

¹For this study, the slowperiod moving average is set to 60 periods, while the fastperiod moving average is set to 30 periods.

exit at each time step, with the position size for each trade is set to 200 USDT.. If the predicted probability exceeds a predefined threshold, the system enforces a liquidation before the traditional trend-reversal signal occurs, aiming to lock in floating profits or mitigate severe drawdowns. By superimposing a data-driven decision layer onto rule-based logic, this architecture significantly improves the precision of position management during holding periods.




Figure 2: Model Architecture

4.1 Feature Engineering

To enable the LSTM decision layer to perceive both the market environment and the current trading state, this paper constructs an input feature system jointly composed of market features and position-context features, providing the informational basis required for the model to identify future unrealized profit drawdown risk. The whole model architecture is shown in Figure 2.

4.1.1 Market Features

When selecting market features, this paper follows the approaches of Fischer and Krauss (2018), Huang et al. (2024), and AlKhatib et al. (2022), and jointly considers price trends, return-generating ability, candlestick patterns, volatility conditions, and trading activity. The detailed definitions of the variables are shown in Table 3.

Table 3: Market Features

Category	Feature	Description
Trend	wma30_rel, wma60_rel	Deviation of the closing price from the moving averages
Trend	wma_spread	Deviation of the 30-day moving average from the 60-day moving average
Trend	wma30_slope, wma60_slope	First-order differences of the moving averages
Trend	wma_spread_slope	The rate of change in the interval between moving averages is used to determine their convergence or divergence.
Trend	wma_spread_accel	Acceleration of movingaverage interval variation
Strength	ADX(14)	Average Directional Index, divided by 100
Momentum	RSI(14)	Relative Strength Index, divided by 100
Momentum	MACD_hist	Moving Average ConvergenceDivergence, with fast period 12, slow period 26 and signal line period 9, normalized by ATR(14)
Return	log_ret, oc_ret	Log return and open-close ratio
Pattern	hl_range, shadow_imbalance	Price range and upper-lower shadow imbalance
Volatility	atr14_ret	ATR(14) divided by the closing price
Trading Activity	vol_z20, taker_buy_ratio	Trading volume Z-score and active buy ratio

4.1.2 Position Context Features

In addition to market features, the key innovation of this paper lies in the introduction of position-context features. Traditional technical indicators usually describe the market itself, but cannot represent the current trading state. However, for the early take-profit problem, the decision implications of the same market trend may differ across different stages of the holding period. For example, the drawdown immediately after opening a position and the drawdown after a large unrealized profit are not equivalent in terms of risk implications. Therefore, this paper represents the position state as model input. The detailed definitions of the variables are shown in Table 4.

Table 4: Position Context Features

Feature	Description
age	Number of bars for which the current position has been held
pnl_dir	Unrealized profit or loss in the direction of the position
mfe	Maximum favorable excursion
retrace	Drawdown from the maximum favorable excursion
dir	Position direction, where “+1” denotes a long position and “-1” denotes a short position

By introducing the above position-context information, the LSTM model no longer makes decisions solely based on market states, but is able to make decisions according to the dynamic profit and loss condition of the current trade, the holding duration, and other information. This allows the model to more closely approximate the risk-management logic of real trading.

4.2 Training Stage

Grid Search: First, we employ a grid search to determine the optimal model parameters. The parameters to be optimized include the drawdown penalty coefficient (λ_{dd}), the rolling window length, the number of neural network layers, and the early-closing decision threshold (θ). We define a discrete search space for these parameters, generating multiple parameter combinations where each corresponds to a distinct model. Subsequently, we train these models individually and evaluate their performance on the validation set to select the top five best-performing models. The scoring mechanism comprehensively considers the Sharpe ratio, the average profit per trade, and the maximum drawdown. We employ the entropy weight method to integrate these three metrics to derive the final evaluation score for the models.

During this process, only models that satisfy the following three strict filtering conditions in the valid set will be retained:

First, the early exit rate must be strictly less than 65% to prevent the model from executing premature exits excessively. Second, the model’s total return must exceed 80% of the total return generated by the baseline strategy, ensuring that the model’s performance does not significantly underperform the baseline. Finally, the maximum drawdown must be less than 150 USDT. Given the previously mentioned position size of 200 USDT per trade, this maximum drawdown corresponds to a maximum loss of 75%. Exceeding this threshold indicates that the strategy has suffered consecutive severe losses within a short period, implying that the model’s exit decisions not only failed to protect

profits but actually exacerbated the losses. The score distribution of the eligible models are displayed in Appendix ?? and the performance of the top five best-performing models on the validation set is presented in Appendix B.2.

Label Construction and Adaptive Margin Setting: To train the LSTM decision layer, a binary supervisory label is constructed for each timestamp during the holding period. Here, Δ_{margin} serves as the core parameter for label definition in this system. A K line is labeled as a positive sample only if the floating return of the current bar exceeds the final return by at least Δ_{margin} . Consequently, the value of Δ_{margin} directly dictates the proportion of positive samples, which in turn influences the exit propensity learned by the model. When the value is too small, numerous marginal samples are labeled as positive, causing the model to frequently trigger premature exit signals. Conversely, an excessively large value leads to an extreme scarcity of positive samples, biasing the model towards non-execution and rendering the early exit mechanism ineffective.

Empirically, constraining the positive sample ratio to approximately 0.35 yields robust model performance. Based on this target ratio, the preliminary Δ_{margin} values computed for each asset are: 0.008 for ETH, 0.006 for BTC, 0.012 for SOL, and 0.008 for XRP. After establishing these initial values, we utilize the top three optimal parameter combinations identified from the preceding grid search to conduct an ablation study. Specifically, we define a set of candidate Δ_{margin} values covering a spectrum from lenient to strict (0.003, 0.005, 0.008, 0.010, 0.012, and 0.015). By exclusively varying Δ_{margin} , we re-evaluate the model’s performance across several key metrics: the positive sample ratio, the validation return versus the baseline strategy return, the Sharpe ratio, and the early exit trigger rate. The cross-asset results of this ablation analysis are detailed in Appendix ??.

To ensure that the selected optimal parameters are not an artifact of overfitting, we conduct a sensitivity analysis on the corresponding parameters for each asset, with the results detailed in Appendix B.3. The findings indicate that the model’s performance does not exhibit high sensitivity to minor variations across the parameters; specifically, no abrupt "isolated peaks" are observed in the plots. The sensitivity analysis for BTC and ETH is particularly robust, suggesting that large-cap cryptocurrencies are relatively insensitive to parameter fluctuations.

Dropout Regularization: To prevent the model from overfitting on the training set, dropout is introduced as a structural regularization technique between the LSTM layers and the fully connected (FC) layer. The initial dropout rate is set to 0.2, serving as a fixed value during the first stage, and is subjected to an ablation study following the completion of the primary grid search.

In terms of experimental design, a two-tier dropout mechanism is implemented at distinct positions. The first tier is the inter-layer dropout, applied between adjacent LSTM layers to randomly mask connections during each forward pass. The second tier is the output dropout, applied between the final hidden state output of the LSTM and

the FC layer, ensuring that the regularization remains effective under both single-layer and multi-layer network configurations.

In the second stage, an ablation study is conducted specifically on the dropout rate. Using the initial dropout rate as an empirical baseline, the optimal parameter combination is fixed after the primary grid search. We then define a candidate set of dropout rates as $\{0.1, 0.2, 0.3\}$ and re-evaluate the model's performance under each setting. The performance evaluation metric is identical to that defined in Equation ??.

4.3 LSTM Model

We adopt a Long Short-Term Memory (LSTM) network as the early-exit decision model. LSTM is suitable for capturing dynamic dependence in financial time series and can, to a certain extent, capture the nonlinear relationship between trend evolution and position-state changes.

The model input is a fixed-length historical feature matrix which we mentioned before:

$$\mathbf{X}_{t-63:t} \in \mathbb{R}^{64 \times 17} \quad (1)$$

where 64 denotes the length of the time window, and 17 denotes the feature dimension contained in each time step.

The LSTM model recursively encodes the input sequence and uses the hidden state of the last time step as a comprehensive representation of the current holding state. The hidden state is then mapped into a one-dimensional output through a fully connected layer and converted into a probability form through the Sigmoid function:

$$p_t = \sigma(\mathbf{w}\mathbf{h}_t + b) \quad (2)$$

where $\mathbf{h}_t$ denotes the hidden state of the LSTM at the current time, $\mathbf{w}$ and b denote the weight and bias of the fully connected layer, respectively, and $\sigma(\cdot)$ denotes the Sigmoid activation function. The final output is interpreted as the probability of early exit at the current time, namely:

$$p_t = P(y_t = 1 \mid \mathbf{X}_{t-63:t}) \quad (3)$$

When p_t is close to 1, it indicates that the model believes the future unrealized profit drawdown risk is relatively high and that the current position should be closed early. When p_t is close to 0, it indicates that continuing to hold the position is more reasonable.

4.4 Decision Rule

In the backtesting stage, the LSTM decision layer operates as an additional exit module on top of the baseline strategy. The baseline strategy still generates trading directions and initial entry signals, while the LSTM model participates in decision-making only when an existing position is being held.

Specifically, when the system executes a WMA crossover rule based on the baseline strategy and establishes a position, for each new candlestick, the market features and position-context features are updated, and the feature sequence of the most recent 64 time steps is input into the LSTM model to obtain the early-exit probability p_t .

The exit rule in this paper can be expressed as:

$$Exit_t = \begin{cases} 1, & p_t > \theta \\ -1, & \text{opposite WMA crossover occurs} \\ 0, & \text{else} \end{cases} \quad (4)$$

where θ is the decision threshold for early exit, determined through grid search. When $p_t > \theta$, it indicates that the model judges the current position to face a relatively high future unrealized profit drawdown risk, and therefore the system will actively close the position before the reverse crossover signal of the baseline strategy appears. If the LSTM does not trigger an early exit, the trade continues to be managed according to the original moving-average reverse crossover rule.

Through the above decision rule, this paper forms a two-layer exit mechanism: the baseline strategy provides trend-reversal exits, while the LSTM layer provides risk-warning-based early exits. The former ensures that the strategy still follows trend-following trading logic, while the latter improves the responsiveness to unrealized profit drawdown risk during the holding period. This mechanism enables the strategy to no longer passively wait for the moving-average reverse crossover, but instead to lock in profits in advance when drawdown risk rises, thereby improving strategy performance.

4.5 Comparison Models

To evaluate the effectiveness of the proposed LSTM-based early-closing framework in enhancing returns and controlling drawdowns, its performance is compared against two representative benchmark strategies. The first is the underlying WMA baseline strategy, which acts as the fundamental trend-following reference utilizing 30-day and 60-day weighted moving averages without any auxiliary exit mechanism. The second is the simple buy-and-hold strategy. This comparative analysis allows us to observe whether the proposed model can significantly outperform the existing benchmark strategies, thereby validating its superiority in dynamic risk identification.

In terms of experimental design, all strategies are uniformly applied to multiple mainstream cryptocurrency trading pairs, specifically BTC/USDT, ETH/USDT, SOL/USDT, and XRP/USDT. Conducting these cross-asset experiments under identical parameters allows for a comprehensive assessment of the robustness and generalizability of the proposed framework across diverse market structures.

5 Results

Based on the validation set performance, we identify the top optimal parameter combinations for each asset (see Table 5). For each parameter combination, we backtest the proposed strategy using the test set data covered from Jul 28, 2025 to Mar 12, 2026 to evaluate the out-of-sample performance of the corresponding model. Furthermore, to benchmark the performance enhancements achieved by our framework, we concurrently evaluate the baseline WMA strategy and the buy-and-hold strategy to analyze the relative improvements.

Table 5: Optimal Hyperparameter Configuration for Each Asset

Asset	Window Size	λ_{dd}	Num Layers	θ	Seed
ETH	96	2.0	1	0.6	42
BTC	32	1.0	2	0.5	456
XRP	64	2.0	2	0.6	123
SOL	32	1.5	1	0.5	123

Note: This table presents the optimal parameter combination for each asset. Specifically, Window Size denotes the lookback window length, λ_{dd} represents the drawdown penalty coefficient, Num Layers indicates the number of neural network layers, θ denotes the early exit decision threshold, and Seed is the random seed.

Figure 3 illustrates the cumulative returns of the combination of the top optimal LSTM early-stopping models, the baseline WMA strategy, and the buy-and-hold strategy for each asset. As can be observed, both the early-closing models and the baseline WMA strategy significantly outperform the buy-and-hold strategy in terms of cumulative returns. With the exception of XRP, the LSTM early-closing models demonstrate a substantial improvement over the baseline strategy across the other three cryptocurrencies.




Figure 3: Total Cumulative Returns of Different Strategies for Each Asset

Note: The solid blue line represents the cumulative returns of the combination of the top three optimal LSTM early-closing models, the dashed line represents the baseline WMA strategy, and the solid orange line represents the buy-and-hold strategy. The shaded area indicates the outperformance (or underperformance) of the early-closing models relative to the baseline strategy.

In addition to cumulative returns, we evaluate the model performance based on the Sharpe ratio and the early-closing probability. Appendix ?? presents the performance metrics of the top three optimal LSTM early-closing models and the baseline WMA strategy for each asset. As can be observed, for every asset, at least one LSTM early-closing model significantly outperforms the baseline strategy. This demonstrates that our proposed model can indeed improve performance on top of the baseline strategy; however, it also indicates that the model’s performance is relatively sensitive to parameter settings.

Furthermore, as previously noted in Section 1, a majority of trades experience a profitable state at some point during the holding period but ultimately close at a loss. To investigate whether the LSTM early-closing model can improve the trade win rate, we evaluate the proportions of losing trades, winning trades, trades converted from losses to profits (relative to the baseline strategy), and trades converted from profits to losses for the top three optimal LSTM early-closing models across each asset, as illustrated in Figure 4 and Table 6.




Figure 4: Win Rate Improvement of Different Strategies for Each Asset

Note: These pie charts illustrate the win rates of the baseline strategy and the top three optimal LSTM early-closing strategies for each asset. Specifically, the blue areas represent trades that remain profitable relative to the baseline strategy, and the gray areas represent trades that remain at a loss. The green areas indicate trades converted from losses to profits compared to the baseline strategy, while the red areas denote trades converted from profits to losses.

Table 6: Win/Loss Transition Summary Across Assets

Asset	Model	Win Rate	Turned Win	Turned Loss
ETH	Baseline	34.7%	—	—
	ws96/ λ 2.0/nl1	38.1%	15 (6.4%)	7 (3.0%)
BTC	Baseline	35.6%	—	—
	ws32/ λ 1.0/nl2	48.1%	34 (14.2%)	4 (1.7%)
XRP	Baseline	38.4%	—	—
	ws64/ λ 2.0/nl2	44.6%	16 (6.9%)	1 (0.4%)
SOL	Baseline	34.9%	—	—
	ws32/ λ 1.5/nl1	40.5%	27 (11.6%)	14 (6.0%)

6 Robustness Checks

Recursive Window Training: In our previous work, the training, validation, and test sets all utilized fixed window lengths. In this study, we employ a recursive-window training method. The initial training set is set to 2020.01–2021.05, the validation set to

2021.05–2021.11, and the test set to 2021.11–2022.06. Thereafter, we extend the sample forward using a 6-month step size, maintaining a consistent ratio of 70%, 15%, and 15% for the training, validation, and test sets within each sample group. The detailed progression is presented in Table 7.

Table 7: Dataset Partitioning Across Expanding Windows

Window	Data Period	Training Set (70%)	Validation Set (15%)	Test Set (15%)
Window 1	2020.01–2022.06	2020.01–2021.05	2021.05–2021.11	2021.11–2022.06
Window 2	2020.01–2023.01	2020.01–2022.03	2022.03–2022.09	2022.09–2023.01
Window 3	2020.01–2023.07	2020.01–2022.09	2022.09–2023.02	2023.02–2023.07
Window 4	2020.01–2024.01	2020.01–2023.03	2023.03–2023.09	2023.09–2024.01
Window 5	2020.01–2026.03	2020.01–2024.04	2024.04–2025.01	2025.01–2026.03

Note: This table presents the temporal boundaries for the expanding data windows. Window 5 represents the original complete dataset. Each window is strictly partitioned chronologically into a training set, a validation set, and a test set to evaluate the out-of-sample robustness of the model across different market periods.

Due to the substantial interference from controversial information regarding XRP and SOL in their early years, we exclude these two assets upon expanding the sample window. The average profitability and Sharpe ratio of the strategies across the test sets of each expanding window for ETH and BTC are illustrated in Figure 5.




Figure 5: Average Profit/Loss of Different Models for ETH and BTC: Expanding Windows

Note: Panels (a) and (b) present the average profitability under different periods for BTC and ETH using expanding windows, respectively. The orange lines depict the Sharpe ratios of the LSTM early-closing models across different configurations. The solid (untextured) bars represent the average profitability of the LSTM early-closing models, while the textured bars denote the average profitability of the baseline WMA strategy.

As can be observed from the figure, for BTC, the LSTM early-closing models underperform the baseline strategy in all but the most recent sample window. In the fifth window, the average profitability of the baseline strategy is -0.9 USDT, whereas the LSTM early-closing model achieves an average profitability of 10.4 USDT. In contrast, for ETH, the LSTM early-closing models consistently outperform the baseline strategy

across the three most recent windows, indicating that our proposed strategy indeed delivers significant performance improvements. Specifically, from the third to the fifth window, the average profitability of the LSTM early-closing models is -7.5 USDT, 227.9 USDT, and 96.8 USDT, respectively, while the corresponding baseline strategy yields only -13.9 USDT, 158.2 USDT, and -12.9 USDT, respectively.

Rolling Window Training: However, the expanding window approach assumes that market conditions remain constant over the long term. In reality, market dynamics are inherently time-varying. Therefore, we employ a rolling window approach to capture these changing market dynamics and to examine whether our proposed strategy can continue to significantly outperform the baseline strategy under this framework. We similarly conducted five sets of tests, and the results are illustrated in Figure 6.



Figure 6: Average Profit/Loss of Different Models for ETH and BTC:Robustness

Note: Panels (a) and (b) present the average profitability under different periods for BTC and ETH using rolling windows, respectively. The orange lines depict the Sharpe ratios of the LSTM early-closing models across different configurations. The solid (untextured) bars represent the average profitability of the LSTM early-closing models, while the textured bars denote the average profitability of the baseline WMA strategy.

Upon switching to the rolling window approach, we find that the overall performance of the LSTM early-closing models on ETH is consistent with that on BTC; for both assets, the early-closing models outperform the baseline strategy in three out of the five test windows. Furthermore, we observe that when the average profitability of the baseline strategy is significantly negative, the LSTM early-closing models can achieve positive returns or incur substantially smaller losses. For instance, for BTC, the baseline strategy incurs losses of -70.0 USDT and -77.7 USDT in the second and fifth windows, respectively, whereas our models limit the losses to merely -15.6 USDT and -30.7 USDT, respectively. This demonstrates that the underlying mechanism by which our proposed strategy improves upon the baseline strategy is not solely locking in profits, but more importantly, executing timely stop-losses.

7 Conclusion and Future Work

The dual-layer framework, which combines a baseline WMA entry rule with an LSTM-based exit policy, significantly outperforms both the standalone baseline strategy and the passive buy-and-hold benchmark. Out-of-sample testing shows substantial improvements in cumulative returns and Sharpe ratios across the majority of tested assets, specifically BTC, ETH, and SOL.

By incorporating transaction-level position-context features specifically maximum favorable excursion (MFE) and trailing drawdown the LSTM model successfully identifies trend exhaustion before a moving average crossover occurs. This early liquidation mechanism effectively addresses the inherent lag of traditional exit rules and protects accumulated paper profits.

Granular per-trade PnL analysis indicates that the LSTM exit system alters the return profile in an asymmetric manner. The framework effectively curtails the downside risks and shifts the median of trade returns upward, while successfully preserving the right-tail gains generated by major macroscopic trends.

While the proposed LSTM exit system delivers robust performance, several avenues remain open for future research. First, while our backtesting accounts for baseline execution costs, the impact of varying market liquidity, slippage, and alternative fee structures under hyper-volatile regimes warrants a more granular microstructure analysis. Second, the current framework operates on an asset-by-asset basis; extending this dual-layer architecture to a multi-asset portfolio level incorporating dynamic capital allocation and cross-asset risk spillovers could further smooth the equity curve. Finally, incorporating Deep Reinforcement Learning (DRL) to jointly optimize entry, sizing, and exit decisions within an end-to-end continuous control environment represents a promising technological advancement over the discrete probability-threshold approach.

References

- Adegboye, A. and Kampouridis, M. (2021). Machine learning classification and regression models for predicting directional changes trend reversal in FX markets. *Expert Systems with Applications*, 173:114645.
- AlKhatib, K., Khazaleh, H., Alkhazaleh, H. A., Al-Soud, A., and Abualigah, L. (2022). A new stock price forecasting method using active deep learning approach. *Journal of Open Innovation: Technology, Market, and Complexity*, 8(2):96.
- Ayala, J., García-Torres, M., Vázquez Noguera, J. L., Gómez-Vela, F., and Divina, F. (2021). Technical analysis strategy optimization using a machine learning approach in stock market indices. *Knowledge-Based Systems*, 225:107119.
- Borovkova, S. and Tsiamas, I. (2019). An ensemble of LSTM neural networks for high-frequency stock market classification. *Journal of Forecasting*, 38(6):600–619.
- Chiarella, C., He, X.-Z., and Hommes, C. (2006). A dynamic analysis of moving average rules. *Journal of Economic Dynamics and Control*, 30(9–10):1729–1753.
- Dai, B., Marshall, B. R., Nguyen, N. H., and Visaltanachoti, N. (2021). Risk reduction using trailing stop-loss rules. *International Review of Finance*, 21(4):1334–1352.
- Dempster, M. A. H. and Jones, C. M. (2001). A real-time adaptive trading system using genetic programming. *Quantitative Finance*, 1(4):397–413.
- Dempster, M. A. H. and Leemans, V. (2006). An automated FX trading system using adaptive reinforcement learning. *Expert Systems with Applications*, 30(3):543–552.
- Ellis, C. A. and Parbery, S. A. (2005). Is smarter better? a comparison of adaptive, and simple moving average trading strategies. *Research in International Business and Finance*, 19(3):399–411.
- Fischer, T. and Krauss, C. (2018). Deep learning with long short-term memory networks for financial market predictions. *European Journal of Operational Research*, 270(2):654–669.
- Gerlein, E. A., McGinnity, M., Belatreche, A., and Coleman, S. (2016). Evaluating machine learning classification for financial trading: An empirical approach. *Expert Systems with Applications*, 54:193–207.
- Ghosh, P., Neufeld, A., and Sahoo, J. K. (2022). Forecasting directional movements of stock prices for intraday trading using LSTM and random forests. *Finance Research Letters*, 46(Part A):102280.

- Glynn, P. W. and Iglehart, D. L. (1995). Trading securities using trailing stops. *Management Science*, 41(6):1096–1106.
- Han, Y., Kim, J., and Enke, D. (2023). A machine learning trading system for the stock market based on N-period Min-Max labeling using XGBoost. *Expert Systems with Applications*, 211:118581.
- Hansun, S. (2013). A new approach of moving average method in time series analysis. In *2013 Conference on New Media Studies (CoNMedia)*, pages 1–4. IEEE.
- Huang, Y., Wan, X., Zhang, L., and Lu, X. (2024). A novel deep reinforcement learning framework with BiLSTM-Attention networks for algorithmic trading. *Expert Systems with Applications*, 240:122581.
- Hunter, J. S. (1986). The exponentially weighted moving average. *Journal of Quality Technology*, 18(4):203–210.
- Hwang, Y., Park, J., Lee, Y., and Lim, D.-Y. (2023). Stop-loss adjusted labels for machine learning-based trading of risky assets. *Finance Research Letters*, 58(Part A):104285.
- Kim, T. and Kim, H. Y. (2019). Optimizing the pairs-trading strategy using deep reinforcement learning with trading and stop-loss boundaries. *Complexity*, 2019:3582516.
- Leung, T. and Li, X. (2015). Optimal mean reversion trading with transaction costs and stop-loss exit. *International Journal of Theoretical and Applied Finance*, 18(3):1–31.
- Leung, T. and Zhang, H. (2021). Optimal trading with a trailing stop. *Applied Mathematics and Optimization*, 83(2):669–698.
- Lo, A. W. and Remorov, A. (2017). Stop-loss strategies with serial correlation, regime switching, and transaction costs. *Journal of Financial Markets*, 34:1–15.
- Marshall, B. R., Nguyen, N. H., and Visaltanachoti, N. (2017). Time series momentum and moving average trading rules. *Quantitative Finance*, 17(3):405–421.
- Mehtab, S., Sen, J., and Dutta, A. (2020). Stock price prediction using machine learning and LSTM-based deep learning models. arXiv preprint arXiv:2009.10819.
- Nelson, D. M. Q., Pereira, A. C. M., and de Oliveira, R. A. (2017). Stock market’s price movement prediction with LSTM neural networks. In *2017 International Joint Conference on Neural Networks (IJCNN)*, pages 1419–1426. IEEE.
- Sebastião, H. and Godinho, P. (2021). Forecasting and trading cryptocurrencies with machine learning under changing market conditions. *Financial Innovation*, 7(1):1–30.

- Shawky, T. (2023). Gold exchange traded funds (ETFs), take profit/stop loss prediction using machine learning. *Journal of Business and Management Studies*, 5(4):6–17.
- Shen, S., Jiang, H., and Zhang, T. (2012). Stock market forecasting using machine learning algorithms. Technical report, Department of Electrical Engineering, Stanford University.
- Sivri, M. S., Gultekin, A. B., Ustundag, A., Beyca, O. F., Ari, E., and Gurcan, O. F. (2023). An intelligent system for determination of stop-loss and take-profit limits: A dynamic decision learning approach. In Kahraman, C., Sari, I. U., Oztaysi, B., Cevik Onar, S., Cebi, S., and Tolga, A. ., editors, *Intelligent and Fuzzy Systems: Intelligence and Sustainable Future Proceedings of the INFUS 2023 Conference*, volume 759 of *Lecture Notes in Networks and Systems*, pages 617–624. Springer Nature.
- Trippi, R. R. and DeSieno, D. (1992). Trading equity index futures with a neural network. *Journal of Portfolio Management*, 19(1):27–33.
- Tsantekidis, A., Passalis, N., Toufa, A.-S., Saitas-Zarkias, K., Chairistanidis, S., and Tefas, A. (2021). Price trailing for financial trading using deep reinforcement learning. *IEEE Transactions on Neural Networks and Learning Systems*, 32(7):2837–2846.
- Vezeris, D., Kyrgos, T., and Schinas, C. (2018). Take profit and stop loss trading strategies comparison in combination with an MACD trading system. *Journal of Risk and Financial Management*, 11(3):56.
- Wang, Z., Huang, B., Tu, S., Zhang, K., and Xu, L. (2021). DeepTrader: A deep reinforcement learning approach for risk-return balanced portfolio management with market conditions embedding. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 643–650.
- Zhong, X. and Enke, D. (2019). Predicting the daily return direction of the stock market using hybrid machine learning algorithms. *Financial Innovation*, 5(1):1–20.

A Additional Tables

B Additional Figures

B.1 Model Scores Across Different Parameter Configurations for Each Asset




Figure 7: Model Scores Across Different Parameter Configurations for ETH

Note: This figure presents the model scores across different parameter configurations for ETH. The orange shaded area indicates the region corresponding to the top five models.




Figure 8: Model Scores Across Different Parameter Configurations for BTC

Note: This figure presents the model scores across different parameter configurations for BTC. The orange shaded area indicates the region corresponding to the top five models.




Figure 9: Model Scores Across Different Parameter Configurations for XRP

Note: This figure presents the model scores across different parameter configurations for XRP. The orange shaded area indicates the region corresponding to the top five models.




Figure 10: Model Scores Across Different Parameter Configurations for SOL

Note: This figure presents the model scores across different parameter configurations for SOL. The orange shaded area indicates the region corresponding to the top five models.

B.2 Validation Set Metrics of Top Five Models for Each Assets




Figure 11: Validation Set Metrics of Top Five Models - ETH

Note: This figure presents the validation set performance on ETH of the top five optimal parameter configurations for the LSTM early-closing model. Specifically, Panel (a) depicts the maximum drawdown of the models, Panel (b) illustrates the average profitability, Panel (c) presents the Sharpe ratio, and Panel (d) shows the early-closing trigger probability.




Figure 12: Validation Set Metrics of Top Five Models - BTC

Note: This figure presents the validation set performance on BTC of the top five optimal parameter configurations for the LSTM early-closing model. Specifically, Panel (a) depicts the maximum drawdown of the models, Panel (b) illustrates the average profitability, Panel (c) presents the Sharpe ratio, and Panel (d) shows the early-closing trigger probability.




Figure 13: Validation Set Metrics of Top Five Models - XRP

Note: This figure presents the validation set performance on XRP of the top five optimal parameter configurations for the LSTM early-closing model. Specifically, Panel (a) depicts the maximum drawdown of the models, Panel (b) illustrates the average profitability, Panel (c) presents the Sharpe ratio, and Panel (d) shows the early-closing trigger probability.




Figure 14: Validation Set Metrics of Top Five Models - SOL

Note: This figure presents the validation set performance on SOL of the top five optimal parameter configurations for the LSTM early-closing model. Specifically, Panel (a) depicts the maximum drawdown of the models, Panel (b) illustrates the average profitability, Panel (c) presents the Sharpe ratio, and Panel (d) shows the early-closing trigger probability.

B.3 Sensitivity Analysis




Figure 15: Sensitivity Analysis for Hyperparameters - ETH

Note: This figure presents a sensitivity analysis of the hyperparameters in the LSTM early-closing model for ETH, including the lookback window, the drawdown penalty coefficient, the number of neural network layers, and the early exit decision threshold θ .

Parameter Sensitivity Analysis — BTC (Top 30% by TOPSIS Score, $n=65$, threshold=0.6637)




Figure 16: Sensitivity Analysis for Hyperparameters - BTC

Note: This figure presents a sensitivity analysis of the hyperparameters in the LSTM early-closing model for BTC, including the lookback window, the drawdown penalty coefficient, the number of neural network layers, and the early exit decision threshold θ .

Parameter Sensitivity Analysis — XRP (Top 30% by TOPSIS Score, $n=28$, threshold=0.5985)




Figure 17: Sensitivity Analysis for Hyperparameters - XRP

Note: This figure presents a sensitivity analysis of the hyperparameters in the LSTM early-closing model for XRP, including the lookback window, the drawdown penalty coefficient, the number of neural network layers, and the early exit decision threshold θ .

Parameter Sensitivity Analysis — SOL (Top 30% by TOPSIS Score, $n=37$, threshold=0.4556)




Figure 18: Sensitivity Analysis for Hyperparameters - SOL

Note: This figure presents a sensitivity analysis of the hyperparameters in the LSTM early-closing model for SOL, including the lookback window, the drawdown penalty coefficient, the number of neural network layers, and the early exit decision threshold θ .